

Using a parallel computer for simple image denoising

Mandatory assignment No. 1 of INF3380

Each student should work independently and submit at Devilry a tar-ball containing the required implementations (to be described below).

1 Introduction

Image denoising refers to the removal of noises from a noise-contaminated image, such that the “smoothed” image more closely resembles the original noise-free image. Since noisy images are present in many real-life situations, image denoising has become an important task in modern use of computers. An example of image denoising is illustrated in Figure 1.




Figure 1: Left: a noisy image; Right: a denoised image.

The purpose of this assignment is to let the students get familiarized with the following important tasks:

1. Translation of simple mathematical formulas to a working code.
2. Compilation of existing C source codes into an external library.
3. Implementation of a simple denoising algorithm.
4. Parallelization of the denoising algorithm via MPI programming.

2 Numerical algorithm

An image can be thought as a 2D array, containing $m \times n$ pixels,

$$\mathbf{u} = \begin{bmatrix} u_{m-1,0} & u_{m-1,1} & \cdots & u_{m-1,n-1} \\ \vdots & \vdots & \vdots & \vdots \\ u_{1,0} & u_{1,1} & \cdots & u_{1,n-1} \\ u_{0,0} & u_{0,1} & \cdots & u_{0,n-1} \end{bmatrix}.$$

A very simple denoising strategy is to apply the so-called *isotropic diffusion* a number of times. During each iteration of isotropic diffusion, we want to compute a “smoothed” version of \mathbf{u} and denote it by $\bar{\mathbf{u}}$. More specifically, we can compute $\bar{u}_{i,j}$ using the following formula:

$$\bar{u}_{i,j} = u_{i,j} + \kappa(u_{i-1,j} + u_{i,j-1} - 4u_{i,j} + u_{i,j+1} + u_{i+1,j}).$$

Here, κ is typically a small constant (such as 0.1). Note that the above formula is used to compute the interior pixels of $\bar{\mathbf{u}}$, that is, $1 \leq i \leq m-2$ and $1 \leq j \leq n-2$. The boundary pixels of $\bar{\mathbf{u}}$ should simply copy the corresponding boundary pixels of \mathbf{u} .

Note. If a new isotropic diffusion iteration is to be carried out, the newly computed pixel values of $\bar{\mathbf{u}}$ should be copied back into \mathbf{u} .

3 Using an external library for reading/writing JPEG images

In order to be able to read in a JPEG image before doing denoising as described above, we will make use of a ready-made external C library package that contains a set of header (*.h) files and C (*.c) files. (You should compile all the *.c files and group the resulting object *.o files into a static library file named libsimplejpeg.a.)

In particular, the following two functions from libsimplejpeg.a will be used later:

```
void import_JPEG_file (const char* filename, unsigned char** image_chars,
                      int* image_height, int* image_width,
                      int* num_components);
void export_JPEG_file (const char* filename, const unsigned char* image_chars,
                      int image_height, int image_width,
                      int num_components, int quality);
```

These two functions can be used, respectively, to read and write a data file of the JPEG format. We remark that each pixel in a grey-scale JPEG image uses one byte, and a one-dimensional array of unsigned char (of total length mn) is used to contain all the pixel data of a grey-scale JPEG image. (In the case of a color JPEG image, a 1D array of `rgbbrgb`... values is read in.)

Moreover, the integer variable `num_components` will contain value 1 after the `import_JPEG_file` function finishes reading a grey-scale JPEG image. Value 1 should also be given to `num_components` before invoking `export_JPEG_file` to export a grey-scale JPEG image. (For a color JPEG image, the value of `num_components` is 3.) We also remark that the last argument `quality` of the `export_JPEG_file` function is an integer indicating the compression level of the resulting JPEG image. A value of 75 is the typical choice of `quality`.

4 Data structure

It should be noted that a 1D array of type `unsigned char` is used by libsimplejpeg.a for reading and writing a JPEG image. A variable of type `unsigned char` always has an integer value between 0 and 255. This is not sufficient for doing accurate denoising computations. To this end, the following data structure can be used to store the $m \times n$ pixel values, in connection with denoising:

```
typedef struct
{
    float** image_data; /* a 2D array of floats */
    int m; /* # pixels in x-direction */
    int n; /* # pixels in y-direction */
}
image;
```

5 Five functions need to be implemented

```
void allocate_image(image *u, int m, int n);
void deallocate_image(image *u);
void convert_jpeg_to_image(const unsigned char* image_chars, image *u);
void convert_image_to_jpeg(const image *u, unsigned char* image_chars);
void iso_diffusion_denoising(image *u, image *u_bar, float kappa, int iters);
```

Function `allocate_image` is supposed to allocate the 2D array `image_data` inside `u`, when `m` and `n` are given as input. The purpose of function `deallocate_image` is to free the storage used by the 2D array `image_data` inside `u`.

Function `convert_jpeg_to_image` is supposed to convert a 1D array of unsigned char values into an image struct. Function `convert_image_to_jpeg` does the conversion in the opposite direction.

The most important function that needs to be implemented is `iso_diffusion_denoising`, which is supposed to carry out `iters` iterations of the isotropic diffusion on a noisy image object `u`. The denoised image is to be stored in the `u_bar` object. Moreover, function `iso_diffusion_denoising` is supposed to be carried out by P MPI processes in collaboration. That is, the `u` and `u_bar` objects on each MPI process actually cover one rectangular region of a large image, which is partitioned and distributed to P processes. In other words, after each iteration of isotropic diffusion, data exchange (using MPI commands) is needed between neighboring MPI processes.

6 Serial implementation

You are required to first write a serial program (named `serial_main.c` that has the following skeleton:

```
#include <stdio.h>
#include <stdlib.h>

// make use of two functions from the simplejpeg library
void import_JPEG_file(const char *filename, unsigned char **image_chars,
                    int *image_height, int *image_width,
                    int *num_components);

void export_JPEG_file(const char *filename, unsigned char *image_chars,
                    int image_height, int image_width,
                    int num_components, int quality);

int main(int argc, char *argv[])
{
    int m, n, c, iters;
    float kappa;
    image u, u_bar;
    unsigned char *image_chars;
    char *input_jpeg_filename, *output_jpeg_filename;

    /* read from command line: kappa, iters, input_jpeg_filename, output_jpeg_filename */
    /* ... */
    import_JPEG_file(input_jpeg_filename, &image_chars, &m, &n, &c);

    allocate_image (&u, m, n);
    allocate_image (&u_bar, m, n);
    convert_jpeg_to_image (image_chars, &u);

    iso_diffusion_denoising (&u, &u_bar, kappa, iters);

    convert_image_to_jpeg (&u_bar, image_chars);
    export_JPEG_file(output_jpeg_filename, image_chars, m, n, c, 75);

    deallocate_image (&u);
    deallocate_image (&u_bar);

    return 0;
}
```

7 Parallel implementation

The skeleton of the parallel program (named `parallel_main.c`) can be as follows:

```
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>

// make use of two functions from the simplejpeg library
void import_JPEG_file(const char *filename, unsigned char **image_chars,
                    int *image_height, int *image_width,
                    int *num_components);

void export_JPEG_file(const char *filename, unsigned char *image_chars,
                    int image_height, int image_width,
                    int num_components, int quality);

int main(int argc, char *argv[])
{
    int m, n, c, iters;
    int my_m, my_n, my_rank, num_procs;
    float kappa;
    image u, u_bar, whole_image;
```

```

unsigned char *image_chars, my_image_chars;
char *input_jpeg_filename, *output_jpeg_filename;

MPI_Init (&argc, &argv);
MPI_Comm_rank (MPI_COMM_WORLD, &my_rank);
MPI_Comm_size (MPI_COMM_WORLD, &num_procs);

/* read from command line: kappa, iters, input_jpeg_filename, output_jpeg_filename */
/* ... */

if (my_rank==0) {
    import_JPEG_file(input_jpeg_filename, &image_chars, &m, &n, &c);
    allocate_image (&whole_image, m, n);
}

MPI_Bcast (&m, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Bcast (&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

/* divide the m x n pixels evenly among the MPI processes */
my_m = ...;
my_n = ...;

allocate_image (&u, my_m, my_n);
allocate_image (&u_bar, my_m, my_n);

/* each process asks process 0 for a partitioned region */
/* of image_chars and copy the values into u */
/* ... */

convert_jpeg_to_image (my_image_chars, &u);
iso_diffusion_denoising (&u, &u_bar, kappa, iters);

/* each process sends its resulting content of u_bar to process 0 */
/* process 0 receives from each process incoming values and */
/* copy them into the designated region of struct whole_image */
/* ... */

if (my_rank==0) {
    convert_image_to_jpeg(&whole_image, image_chars);
    export_JPEG_file(output_jpeg_filename, image_chars, m, n, c, 75);
    deallocate_image (&whole_image);
}

deallocate_image (&u);
deallocate_image (&u_bar);

MPI_Finalize ();
return 0;
}

```

8 A test case

As an example, a noisy jpeg picture of Mona Lisa (available from the INF3380 course website) can be used to test your serial and parallel implementation. Your implementation is expected to follow the above skeletons. A gzipped tar-ball named `oblig1template.tar.gz` (also available from the course website) can be used as the template.